

GAME OF CLASSIFICATION

Deep Learning on Game of Thrones Screenshots

Machine Learning II
Group 5
Michael Siebel


PURPOSE

- Game of Thrones was HBO's most popular TV series
 - It made HBO substantial profits
 - But, it was also the most pirated show in history
- My goal:
 - To create a copyright detection algorithm
 - Classify screenshots of Game of Thrones against a variety of other TV shows from another genre

A common issue television networks face is improper distribution of their intellectual property

CHALLENGE

- Game of Thrones is known for having a very high number of characters & sets
 - The model should not try to classify faces
 - Instead, general aesthetic
- Attributes of importance:
 - High production values
 - Color treated
 - Costume drama
 - Outdoor/rustic castles sets



COMPARISON GROUP

- All other shows are sitcoms
 - 13 sitcoms, *1989-Today*
 - Again, many characters & sets
 - Has sitcom aesthetic
- Attributes
 - Lower production values
 - Colorful/not color treated
 - Modern clothes/locations
 - Indoor sets



DATA

- Webscrapped all photos from the Internet Movie Database (IMDB.com)
 - All are screenshots (no BTS or promos)
 - Images were thumbnails
 - Only 100x100 pixels
- Pre-cropped photos
 - Focus more on faces
 - Backgrounds can be hidden
- Somewhat imbalanced data

Images

Total: 13,535

Game of Thrones: 3,883

Sitcom: 9,653

*Game of Thrones comprises
~30% total data*

MODELING

- Modeling with a Convolutional Neural Network (CNN)
 - Key attributes relate to RGB values
 - Image filters capture general aesthetic better than Multilayer Perceptron (MLP)
- Solving binary classification problem
 - Splitting data into training (70%), testing (15%), and validation (15%)
 - Training on a Binary Cross Entropy loss (BCELoss) function
 - Evaluating testing data on several models based on F1 scores
 - Evaluating final validation models based on F1 scores
- Coding is all done in Pytorch

DATA AUGMENTATION

- Oversampled Game of Thrones
 - Only on the training set
 - Doubled images
 - Shuffled into main training
- More balanced
 - Went from 29% to 44% of images

Labels	0 Sitcom	1 Game of Thrones
Pre-Oversampling	6763	2711
Post-Oversampling	6763	5422

```
# Oversample GoT
x_ovsp = []
y_ovsp = []
for im, lab in zip(x_train, y_train):
    if lab == 1:
        x_ovsp.append(im)
        y_ovsp.append(lab)

y_ovsp = np.array(y_ovsp)

def shuffle_train(a, b):
    assert len(a) == len(b)
    shuffled_a = np.empty(a.shape, dtype=a.dtype)
    shuffled_b = np.empty(b.shape, dtype=b.dtype)
    permutation = np.random.permutation(len(a))
    for old_index, new_index in enumerate(permutation):
        shuffled_a[new_index] = a[old_index]
        shuffled_b[new_index] = b[old_index]
    return shuffled_a, shuffled_b

x_train = np.concatenate((x_train, x_ovsp), axis=0)
y_train = np.concatenate((y_train, y_ovsp), axis=0)
x_train, y_train = shuffle_train(x_train, y_train)
```

DATA AUGMENTATION

- Resized images to half their size
- On training:
 - Used proportional color jitter
 - Random flips and rotation
 - Normalized
- On testing:
 - Center crop
 - Normalized
 - Avoided randomized transformations


```
# Data augmentation on photos
resize = 50

tf_train = transforms.Compose([
    transforms.RandomResizedCrop(resize),
    transforms.ColorJitter(.3, .3, .3),
    transforms.RandomHorizontalFlip(),
    transforms.RandomRotation(20),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])

tf_test = transforms.Compose([
    transforms.CenterCrop(resize),
    transforms.ToTensor(),
    transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
])
```


PARAMETERS AND NETWORK ARCHITECTURE


- Began with 5 convolutional layers
 - Batch normalization on each layer
 - 0.2 dropouts after layer 3, layer 5, and fully connected layer
- Oversampled (doubled) Game of Thrones
- Used only 25 epochs
- Evaluated best learning rate based on F1 score on test dataset



PARAMETERS AND NETWORK ARCHITECTURE

- Updated epochs to 300
 - Best learning rate became smaller
 - F1 scores did not necessarily improve
- Over predicting Game of Thrones
 - Oversampling is working too well

Confusion Matrix	Predicted Sitcom	Predicted Game of Thrones
Actual Sitcom	701	725
Actual Game of Thrones	290	314



FINAL MODEL ARCHITECTURE

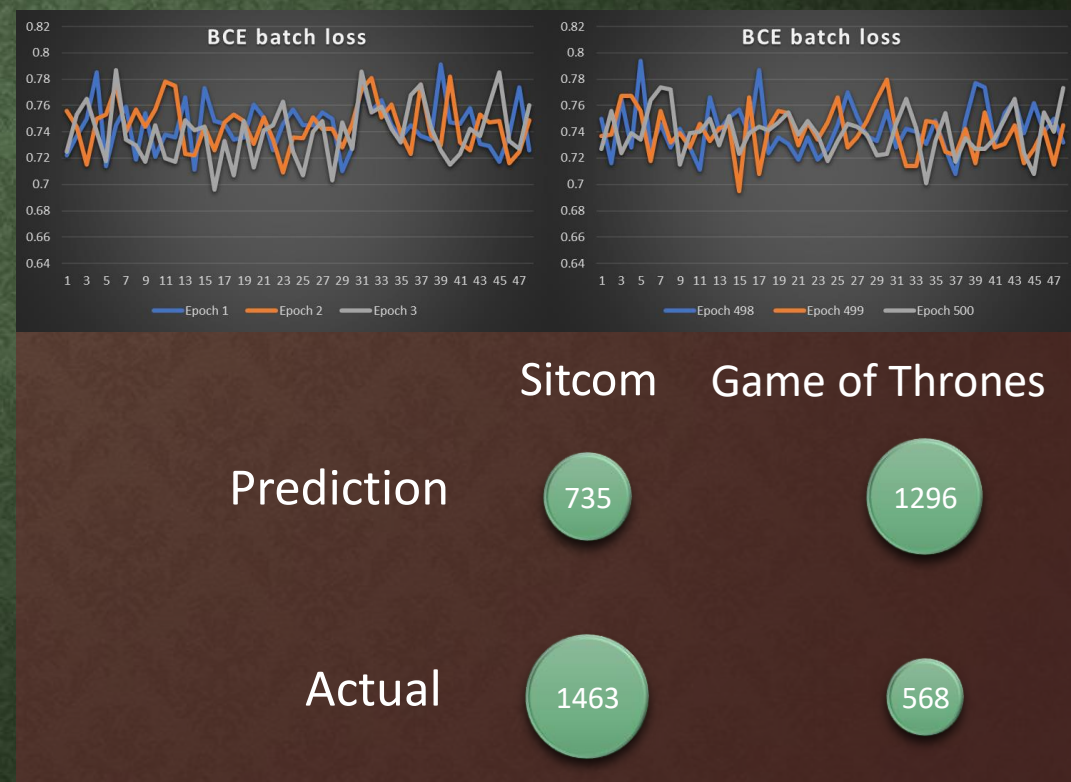
- Updated epochs to 500
- Updated learning rate to $1e-8$
 - After attempting a learning rate of $1e-7$
 - Received similar F1-score but more true positives with the lower rate
- Increased dropouts to 5x

	Input	Output	Activation	Batch Normalization	Pooling	Dropout
Convolutional Layer 1	3	16	ReLU	2D	Max Pooling	
Convolutional Layer 2	16	32	ReLU	2D	Max Pooling	0.2
Convolutional Layer 3	32	64	ReLU	2D	Max Pooling	0.2
Convolutional Layer 4	64	128	ReLU	2D	Max Pooling	0.2
Convolutional Layer 5	128	256	ReLU	2D	Max Pooling	0.2
Fully Connected Layer	256	128	ReLU	1D		0.2
Output layer	128	1	Sigmoid			

MAIN MODEL PERFORMANCE

- Ran on validation dataset
 - Achieved F1-score of 0.39
 - Achieved accuracy of 0.44
- Greatly over predicted
- Loss function did not improve over epochs

Confusion Matrix	Predicted Sitcom	Predicted Game of Thrones
Actual Sitcom	533	930
Actual Game of Thrones	202	366



ALTERNATIVE MODELS

- Experimented on the main model by
 - Removing oversampling
 - Reducing to 3 convolutional layers
 - Running a MLP version
- Surprisingly, the MLP version had similar classification scores as the main model
 - Achieved same F1-Score
 - Received slightly less true positives (342)

No Oversampling

F1-Score 0.29

Accuracy 0.60



Three Layers

F1-Score 0.35

Accuracy 0.53



MLP Version

F1-Score 0.39

Accuracy 0.47

CONCLUSION

- Learning rate was the most tuned hyper-parameter

Did not do enough

- Epochs helped very little

Allowed lower learning rates

Did not lower the loss metrics much

- Oversampling only slightly helped

Improved precision, lowered recall

- CNN surprisingly was not much better than MLP

Image size was manageable by MLP

Full sized versions would not be

- Layers make a difference

Three layers performed much worse

Small images, though, can only have so many pooling layers

NEXT STEPS

- Image quality was likely too weak for the complex analysis of which CNN's are capable

Would need to get full-size images for better analysis

- The loss function may be at issue

Need to better investigate why the loss metrics did not diminish greatly

- Ensemble algorithms are likely necessary
- ...Or simply try classifying the Simpsons